

the Semantic SQL Transducer

seamless knowledge graphs and databases interoperation

Enrico Franconi, KRDB - Free University of Bozen-Bolzano



Agenda

The Semantic SQL Transducer *FOR THE* RDF CROWD

- Start from a relational database
- Find its core (i.e., lossless) conceptual schema
- Express the core conceptual schema in RDFS + SHACL
- Materialise the relational database in the core conceptual schema, maintaining the bidirectional connection
- The relational database can now be queried and updated via SPARQL (and dually the RDF graph can be queried and updated via SQL)
- Non-trivial technical issues, but the transducer is generated automatically from the database
- Implemented in SQL

ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

ssn, phone, email \rightarrow name, dept, manager

ssn \rightarrow name, dept

ssn \twoheadrightarrow phone

ssn \twoheadrightarrow email

dept \rightarrow manager

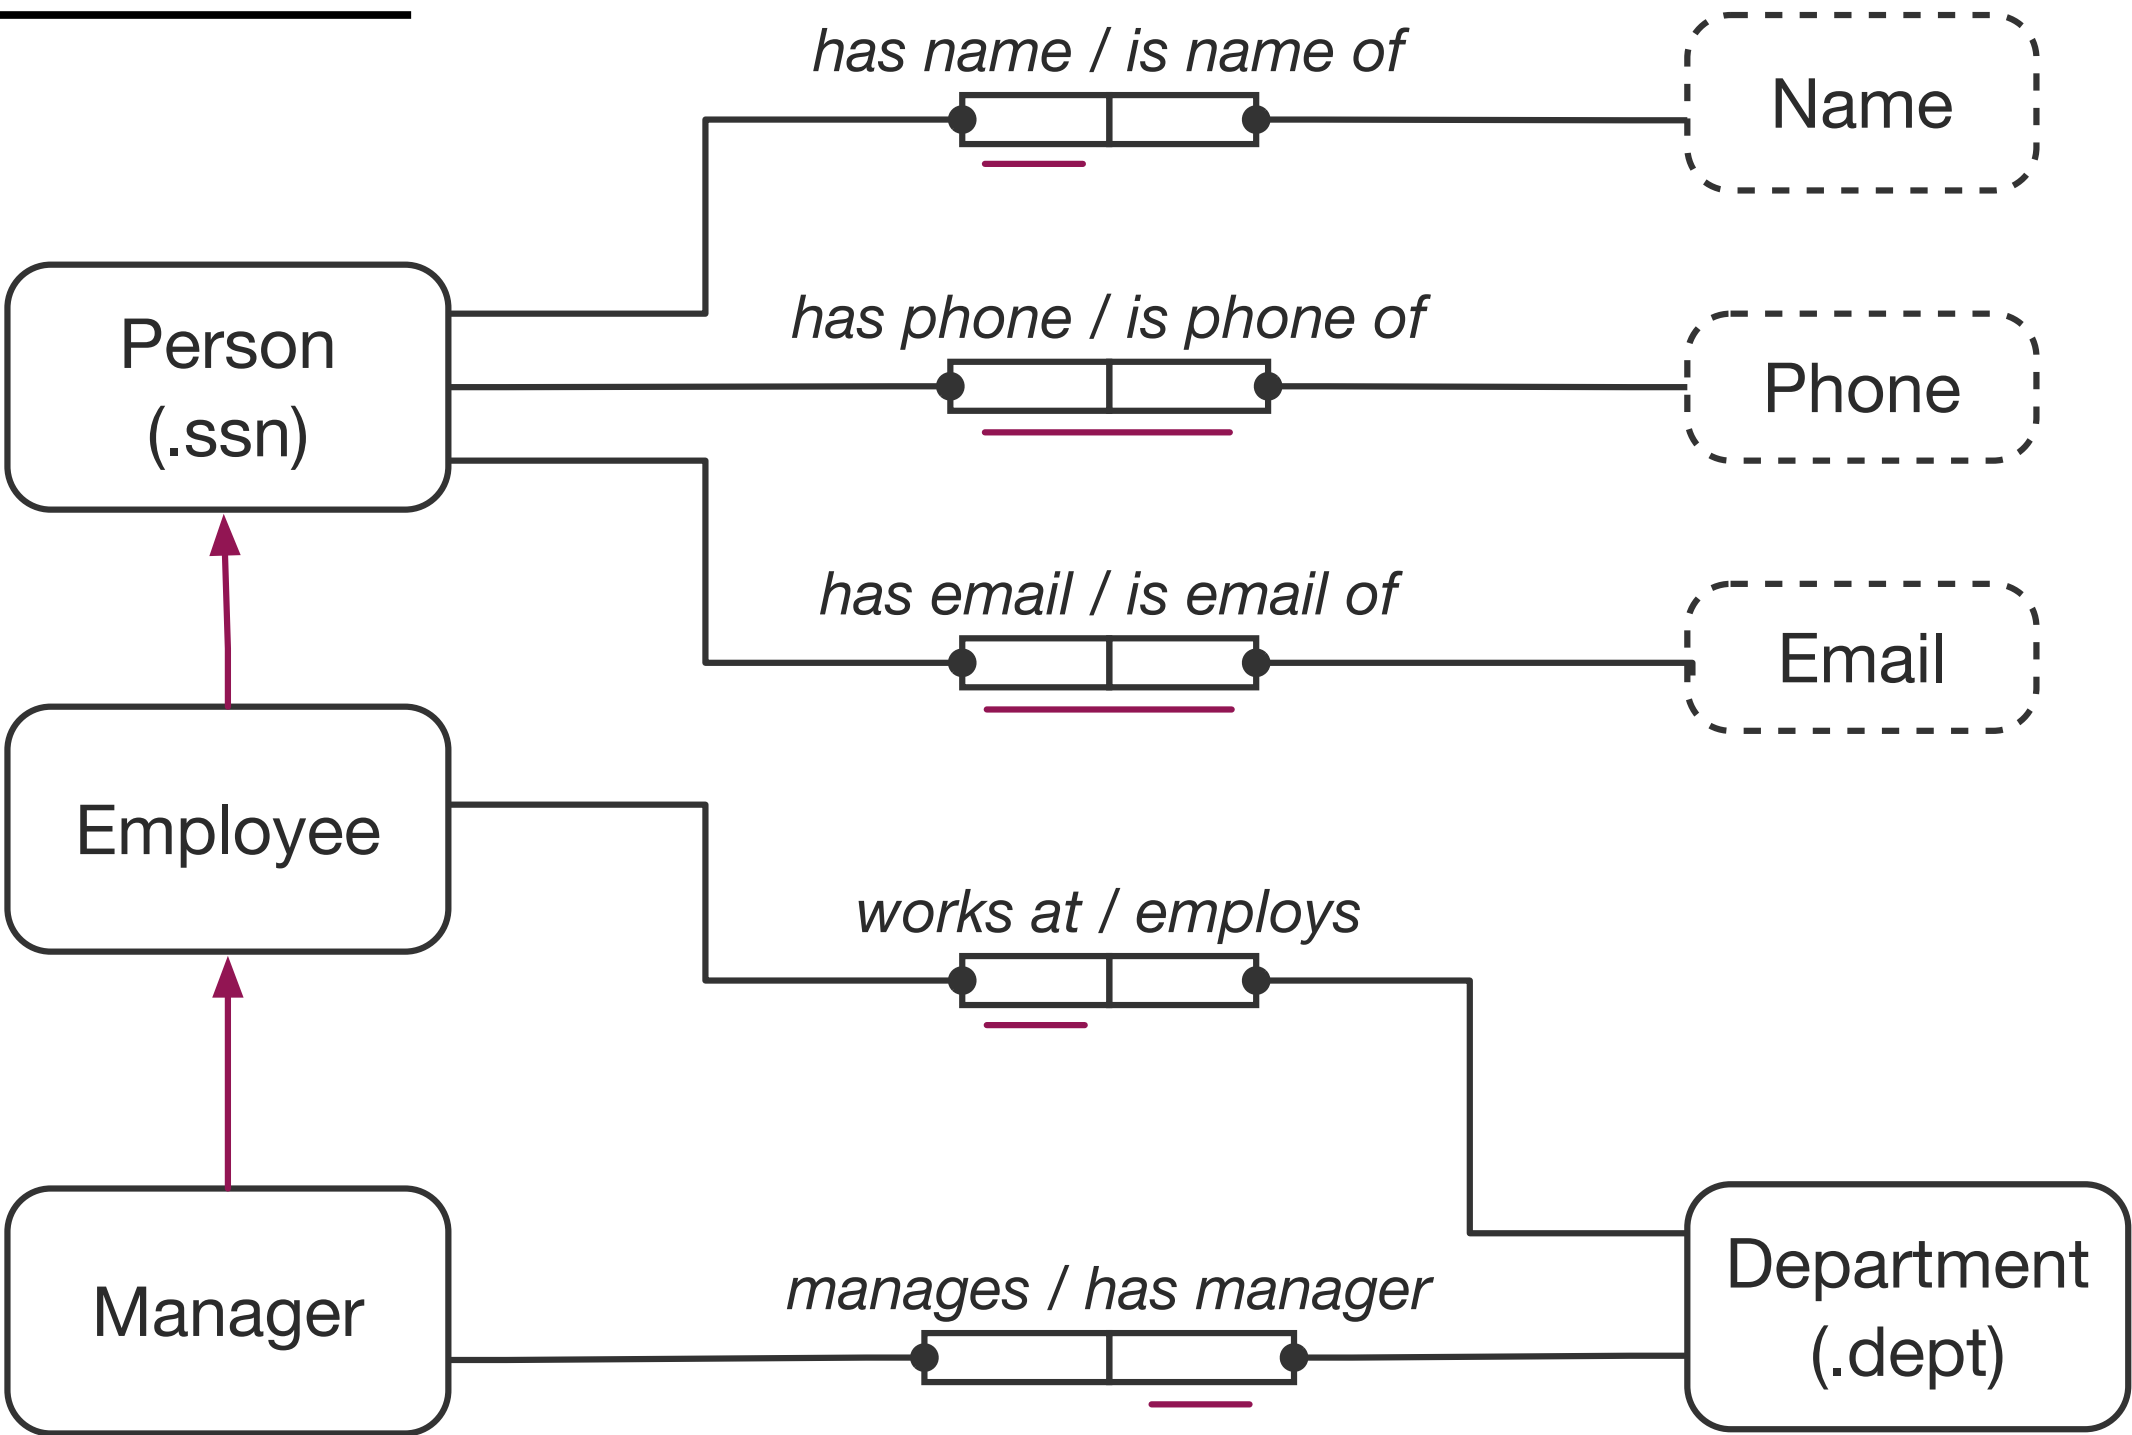
jointly_nullable(dept, manager)

$\sigma_{\text{not_null}(\text{dept})}(\text{manager} \subseteq \text{ssn})$

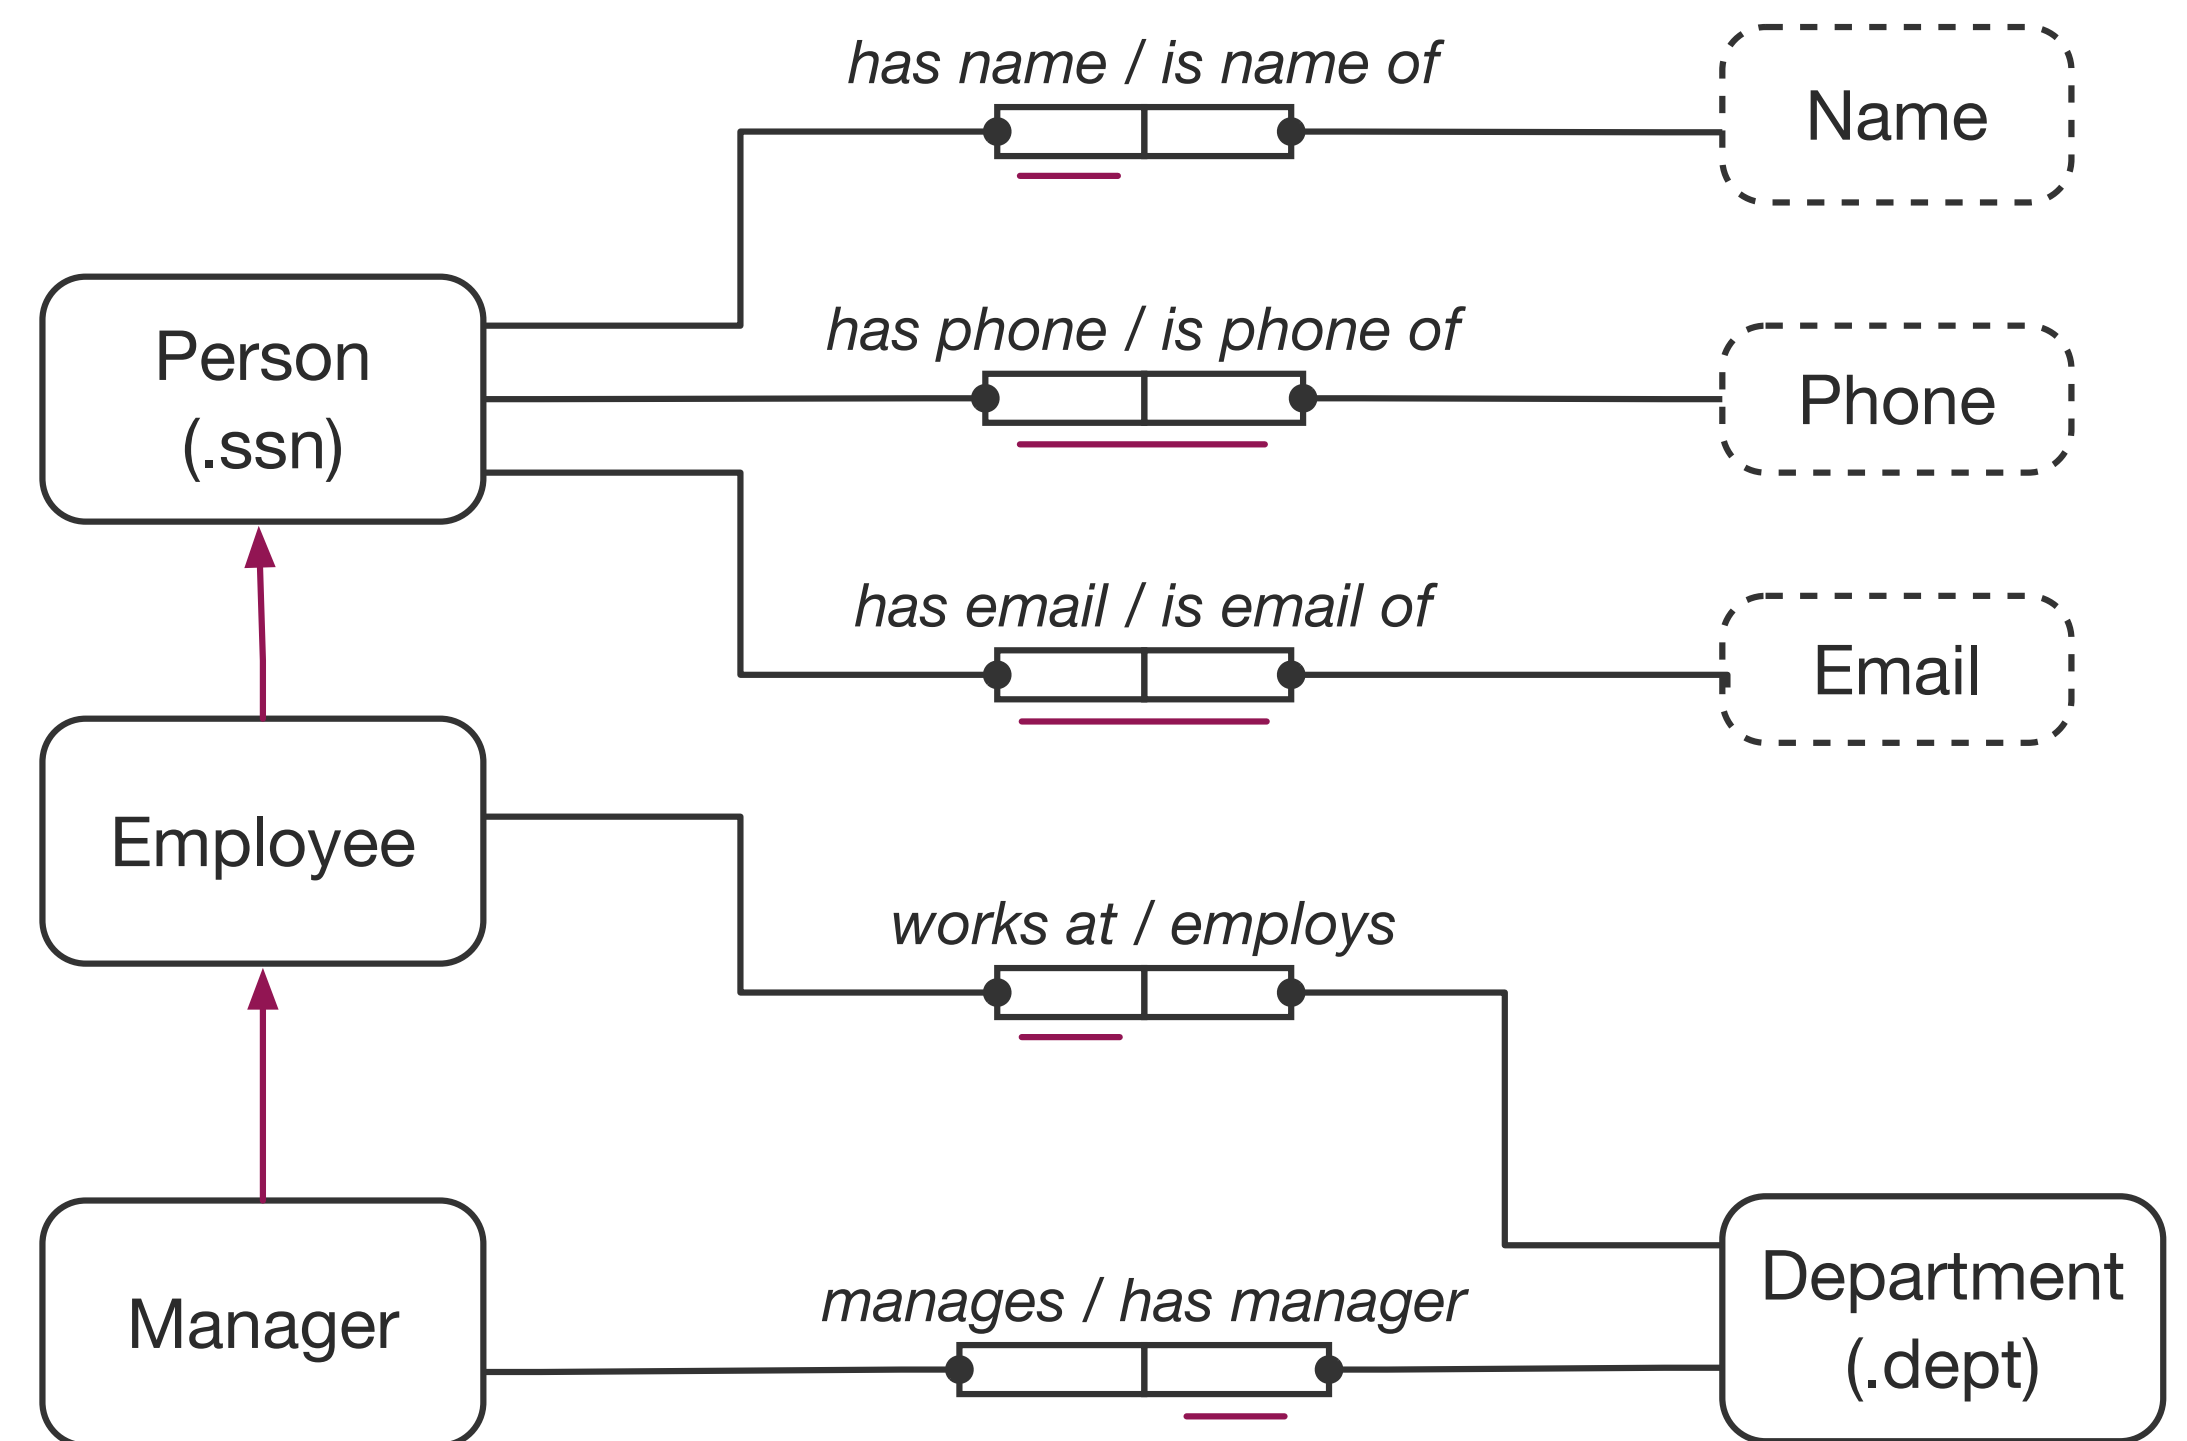
ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

THE CORE CONCEPTUAL SCHEMA

ssn, phone, email → name, dept, manager
ssn → name, dept
ssn → phone
ssn → email
dept → manager
jointly_nullable(dept, manager)
σ_{not_null(dept)} (manager ⊆ ssn)



THE CORE CONCEPTUAL SCHEMA



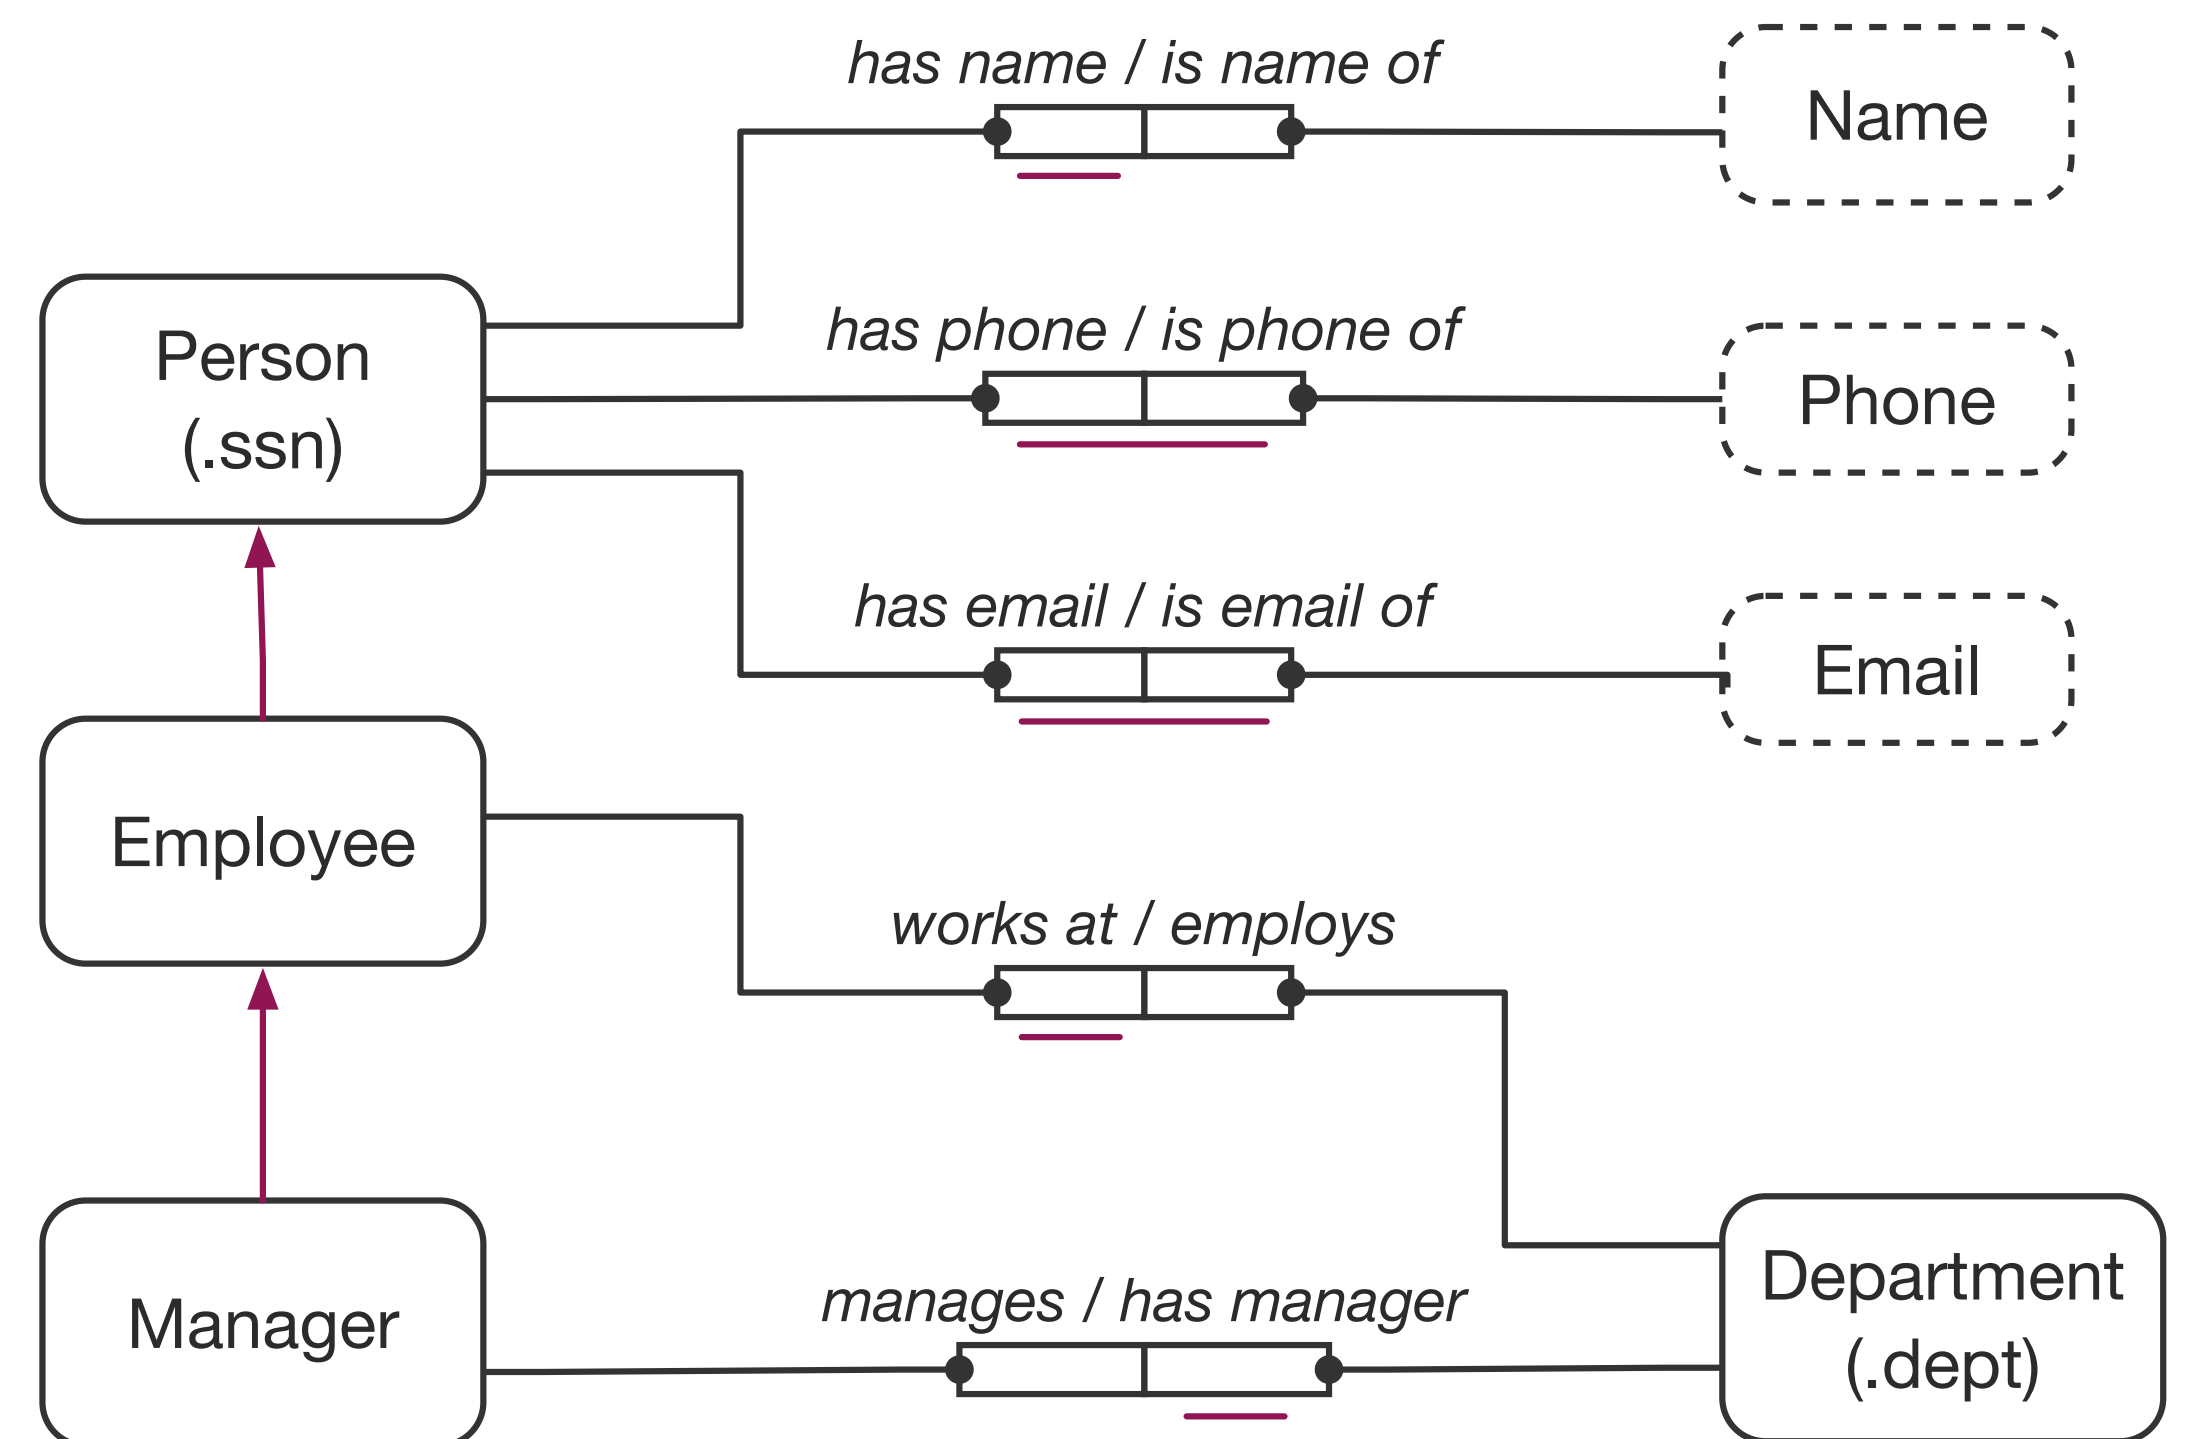
:Manager **rdfs:subClassOf** :Employee .
:Employee **rdfs:subClassOf** :Person .

:has-ssn **rdfs:domain** :Person .
:has-ssn **rdfs:range** xsd:string .
:has-dept **rdfs:domain** :Department .
:has-dept **rdfs:range** xsd:string .

:has-name **rdfs:domain** :Person .
:has-name **rdfs:range** xsd:string .
:has-phone **rdfs:domain** :Person .
:has-phone **rdfs:range** xsd:string .
:has-email **rdfs:domain** :Person .
:has-email **rdfs:range** xsd:string .

:employs **rdfs:domain** :Department .
:employs **rdfs:range** :Employee .
:manages **rdfs:domain** :Manager .
:manages **rdfs:range** :Department .

+ cardinalities:



ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

:Person/ssn=ssn1 **rdf:type** :Person .
 :Person/ssn=ssn1 **:has-ssn** "ssn1" .
 :Person/ssn=ssn2 **rdf:type** :Employee .
 :Employee/ssn=ssn2 **:has-ssn** "ssn2" .
 :Employee/ssn=ssn2 **:has-name** "Mary" .
 :Employee/ssn=ssn2 **:has-phone** 2345 .
 :Employee/ssn=ssn2 **:has-phone** 3456 .
 :Department/dept=D1 **rdf:type** :Department .
 :Department/dept=D1 **:employs** :Employee/ssn=ssn2 .
 :Employee/ssn=ssn3 **rdf:type** :Manager .
 :Employee/ssn=ssn3 **:manages** :Department/dept=D1 .
 ...etc...etc...etc...etc...etc...etc...

+

:Manager **rdfs:subClassOf** :Employee .
 :Employee **rdfs:subClassOf** :Person .

 :has-ssn **rdfs:domain** :Person .
 :has-ssn **rdfs:range** xsd:string .
 :has-dept **rdfs:domain** :Department .
 :has-dept **rdfs:range** xsd:string .

 :has-name **rdfs:domain** :Person .
 :has-name **rdfs:range** xsd:string .
 :has-phone **rdfs:domain** :Person .
 :has-phone **rdfs:range** xsd:string .
 :has-email **rdfs:domain** :Person .
 :has-email **rdfs:range** xsd:string .

 :employs **rdfs:domain** :Department .
 :employs **rdfs:range** :Employee .
 :manages **rdfs:domain** :Manager .
 :manages **rdfs:range** :Department .

+ cardinalities:




ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

INSERT DATA

```
{ :Person/ssn=ssn1 :has-phone 6789 . }
```

INSERT DATA

```
{ :Person/ssn=ssn1 :has-name 'John Doe' . }
```

INSERT DATA

```
{  
:Person/ssn=ssn5 :has-name 'Linda' .  
:Person/ssn=ssn5 :has-ssn 'ssn5' .  
:Person/ssn=ssn5 :has-phone 1234 .  
:Person/ssn=ssn5 :has-email 'linda@company.com'  
:Department/dept=2 :employs :Person/ssn=ssn5 .  
}
```

DELETE DATA

```
{ :Person/ssn=ssn2 :has-phone 2345 . }
```

DELETE DATA

```
{ :Person/ssn=ssn4 :has-phone 5678 . }
```

```
INSERT INTO table
VALUES ('ssn1', 'John', 1234, 'john@company.com', NULL, NULL);
```

ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

```
INSERT INTO table
VALUES ('ssn5', 'Linda', 1234, 'linda@company.com', 'D2', 'ssn4');
```

```
DELETE FROM table
WHERE ssn = 'ssn4' AND phone = 5678 ;
```


```
INSERT DATA
{ :Person/ssn=ssn1 :has-phone 6789 . }
```

```
INSERT DATA
{ :Person/ssn=ssn1 :has-name 'John Doe' . }
```

```
INSERT DATA
{
:Person/ssn=ssn5 :has-name 'Linda' .
:Person/ssn=ssn5 :has-ssn 'ssn5' .
:Person/ssn=ssn5 :has-phone 1234 .
:Person/ssn=ssn5 :has-email 'linda@company.com'
:Department/dept=2 :employs :Person/ssn=ssn5 .
}
```

```
DELETE DATA
{ :Person/ssn=ssn2 :has-phone 2345 . }
```

```
DELETE DATA
{ :Person/ssn=ssn4 :has-phone 5678 . }
```



LOSSLESS DYNAMIC REPLICATION

ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

+ *constraints*

Person

PID

' :Person/ssn=ssn1'
' :Person/ssn=ssn2'
' :Person/ssn=ssn3'
' :Person/ssn=ssn4'

Employee

PID

' :Person/ssn=ssn2'
' :Person/ssn=ssn3'
' :Person/ssn=ssn4'

Manager

PID

' :Person/ssn=ssn3'
' :Person/ssn=ssn4'

Department

DID

' :Department/dept=D1'
' :Department/dept=D2'

has-ssn

PID

PID	ssn
' :Person/ssn=ssn1'	'ssn1'
' :Person/ssn=ssn2'	'ssn2'
' :Person/ssn=ssn3'	'ssn3'
' :Person/ssn=ssn4'	'ssn4'

has-email

PID

PID	email
' :Person/ssn=ssn1'	'john@company.com'
' :Person/ssn=ssn2'	'mary@company.com'
' :Person/ssn=ssn2'	'mary@gmail.com'
' :Person/ssn=ssn3'	'sue@company.com'
' :Person/ssn=ssn4'	'paul@company.com'


employs

DID

DID	PID
' :Department/dept=D1'	' :Person/ssn=ssn3'
' :Department/dept=D2'	' :Person/ssn=ssn4'

...etc...etc...etc...etc...etc...etc...

+ *constraints*



ssn	name	phone	email	dept	manager
'ssn1'	'John'	1234	'john@company.com'	NULL	NULL
'ssn2'	'Mary'	2345	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@company.com'	'D1'	'ssn3'
'ssn2'	'Mary'	2345	'mary@gmail.com'	'D1'	'ssn3'
'ssn2'	'Mary'	3456	'mary@gmail.com'	'D1'	'ssn3'
'ssn3'	'Sue'	4567	'sue@company.com'	'D1'	'ssn3'
'ssn4'	'Paul'	5678	'paul@company.com'	'D2'	'ssn4'

+ constraints

<i>Person</i>	<i>Employee</i>	<i>Manager</i>	<i>Department</i>
PID	PID	PID	DID
':Person/ssn=ssn1'	':Person/ssn=ssn2'	':Person/ssn=ssn3'	':Department/dept=D1'
':Person/ssn=ssn2'	':Person/ssn=ssn3'	':Person/ssn=ssn4'	':Department/dept=D2'
':Person/ssn=ssn3'			
':Person/ssn=ssn4'			

<i>has-ssn</i>	<i>has-email</i>
PID	PID
ssn	email
':Person/ssn=ssn1'	':Person/ssn=ssn1'
':Person/ssn=ssn2'	':Person/ssn=ssn2'
':Person/ssn=ssn3'	':Person/ssn=ssn2'
':Person/ssn=ssn4'	':Person/ssn=ssn3'
	':Person/ssn=ssn4'

<i>employs</i>	...
DID	PID
':Department/dept=D1'	':Person/ssn=ssn3'
':Department/dept=D2'	':Person/ssn=ssn4'

...etc...etc...etc...etc...etc...etc...
+ constraints

Advantages

- native **queries** possible at each end, with local response time
- native **update** possible at each end, with local response time, with **propagation** of the update at the other end, and **integrity** maintained at both ends
- supports native ACID **transactions**
- given a source database in a RDBMS, the transducer is just a bunch of SQL **materialised views with constraints** and SQL **triggers** added to the RDBMS, without changing the source database
- given a source database in a RDBMS, the transducer is generated by an **automatic** procedure, which applies lossless transformation patterns
- the designer **interacts** by providing class names and predicate names
- implemented in **PostgreSQL**

Advantages, Disadvantages

- native **queries** possible at each end, with local response time
- native **update** possible at each end, with local response time, **with propagation of the update at the other end**, and **integrity** maintained at both ends
- supports native ACID **transactions**
- given a source database in a RDBMS, the transducer is just a bunch of SQL **materialised views with constraints** and SQL **triggers** added to the RDBMS, without changing the source database
- given a source database in a RDBMS, the transducer is generated by an **automatic** procedure, which applies lossless transformation patterns
- the designer **interacts** by providing class names and predicate names
- implemented in **PostgreSQL**

Advantages, Disadvantages, Status

- native **queries** possible at each end, with local response time
- native **update** possible at each end, with local response time, **with propagation of the update at the other end**, and **integrity** maintained at both ends
- supports native ACID **transactions**
- given a source database in a RDBMS, the transducer is just a bunch of SQL **materialised views with constraints** and SQL **triggers** added to the RDBMS, without changing the source database
- given a source database in a RDBMS, the transducer is generated by an **automatic** procedure, which applies lossless transformation patterns
- the designer **interacts** by providing class names and predicate names
- implemented in **PostgreSQL**

And now?

(seeking collaborations)

- **Open-source serious implementation**
- Documentation of the methodology
- Also:
 - collaboration with RDBMS profiling companies
 - deployment in enterprise settings